

Introduction to Kubernetes/MiniKube (Lightning talk and MiniKube Demo)

By George Franklin

Our Main Sponsors



Limerick AI – (Applications and Games) Software Development MeetUp

<http://meetup.com/LimerickAiSD>

Twitter @LimerickAiSD

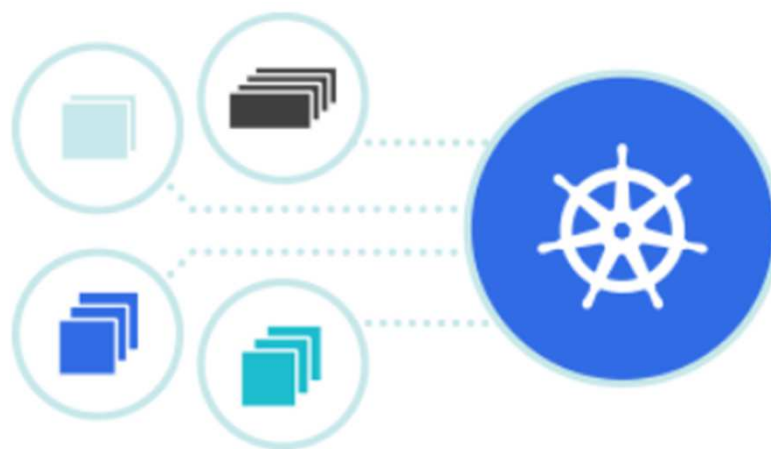
Big Thanks To

- For Everyone Joining & Attending
- A Big thanks to Bank of Ireland for the BOI workbench meetup location in Limerick.
- Our Sponsors JetBrains and Redgate for their raffle product licenses.

What is Kubernetes

So what will we cover

- What is Kubernetes
 - What are containers / pod
 - Why do containers need orchestration
 - The main working processes/components of kubernetes
 - It all about Pods and not Containers !!
-
- How we can control a cluster via cloud tools, the Apiserver and the client application – Kubectl.
 - How to start a statefulset and a simple pod using Kubectl
-
- We won't have time to cover all the Controller types, that's for another talk perhaps with a focus on building microservices or games services in kubernetes.



Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.

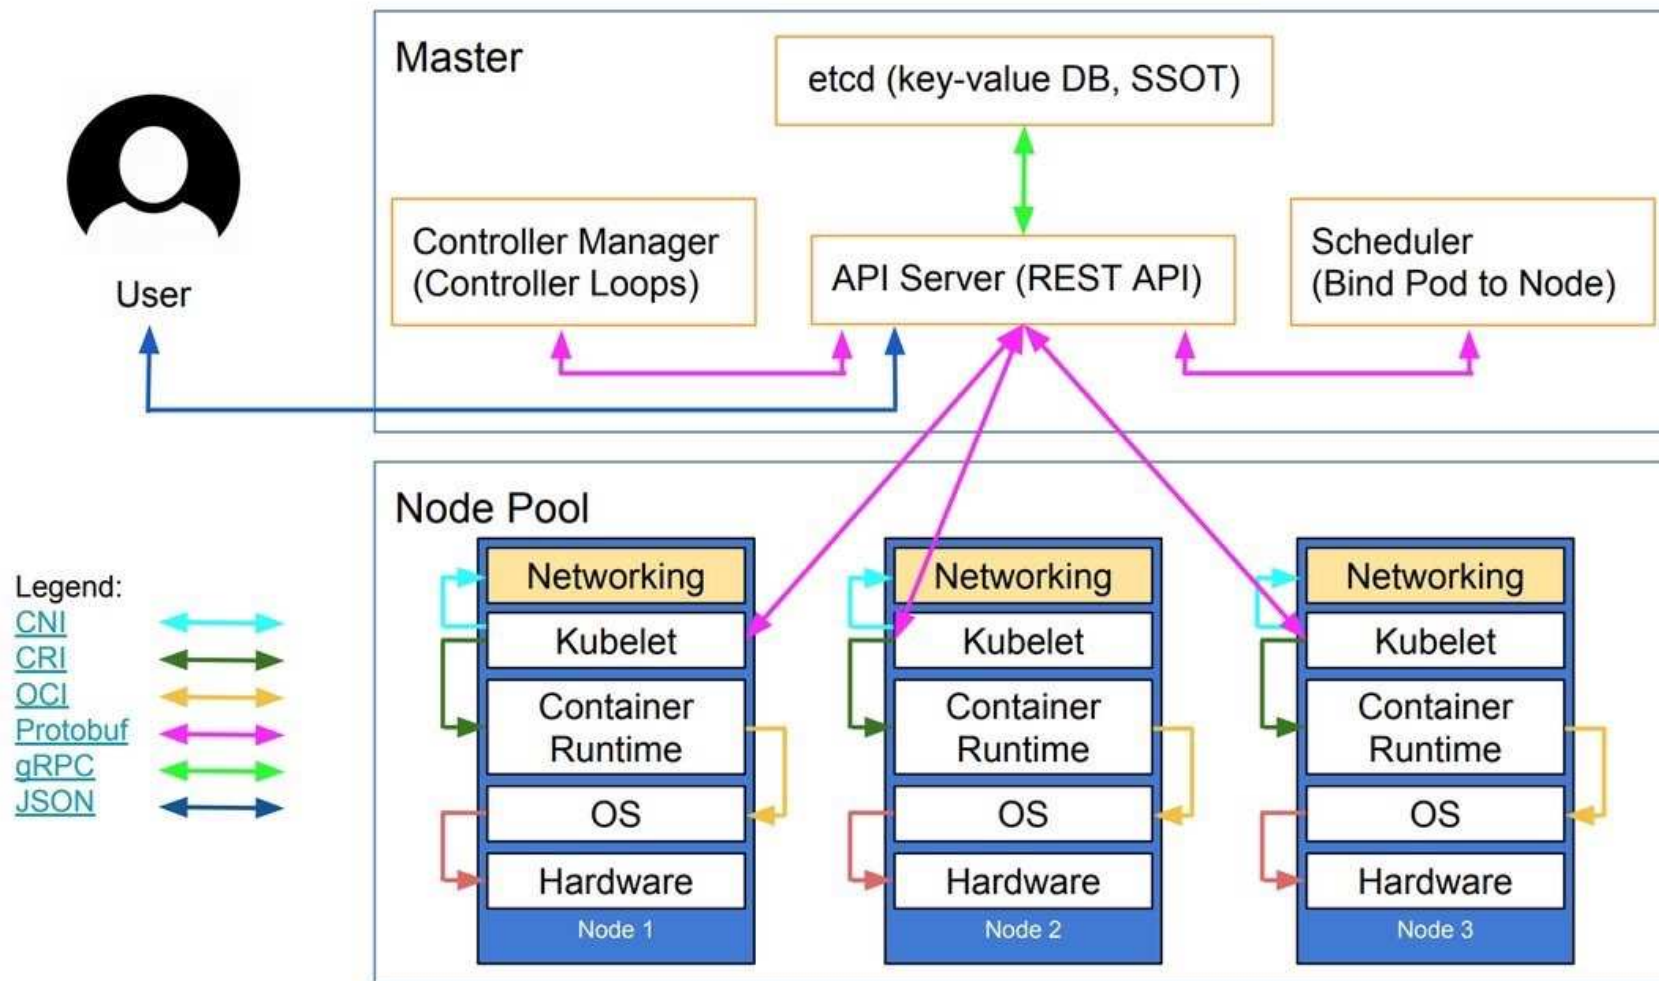
It groups containers that make up an application into logical units for easy management and discovery. Kubernetes builds upon [15 years of experience of running production workloads at Google](#), combined with best-of-breed ideas and practices from the community.



Kubernetes Cluster

- We define **Kubernetes** as a **container orchestration framework**.
- A **Kubernetes Cluster** is used to describe the **deployment of a master and worker nodes**. Users **application containers will run on the worker nodes**.
- **Kubernetes** is also a **Hybrid system** as part of the framework uses both containers as well as system processes to actually **manage the cluster spread across the master and worker nodes**. So you could say it eats its own dog food.
- So we mentioned **containers** and specifically lets look at **Docker** the most popular container technology at the moment used by default in **Kubernetes**. (But you can use others i.e. **rkt**)
- Also ask the question - **Why do containers need orchestration** and the most common **uses cases of large deployments of containers**.

Kubernetes' high-level component architecture



Picture Source : <http://sysadvent.blogspot.com/2017/12/day-24-on-premise-kubernetes-with.html> You can use the picture below for visualisation. Thanks to Lucas Kälström for creating this (@kubernetesonarm), used in his presentation on KubeCon

Kubernetes and Docker

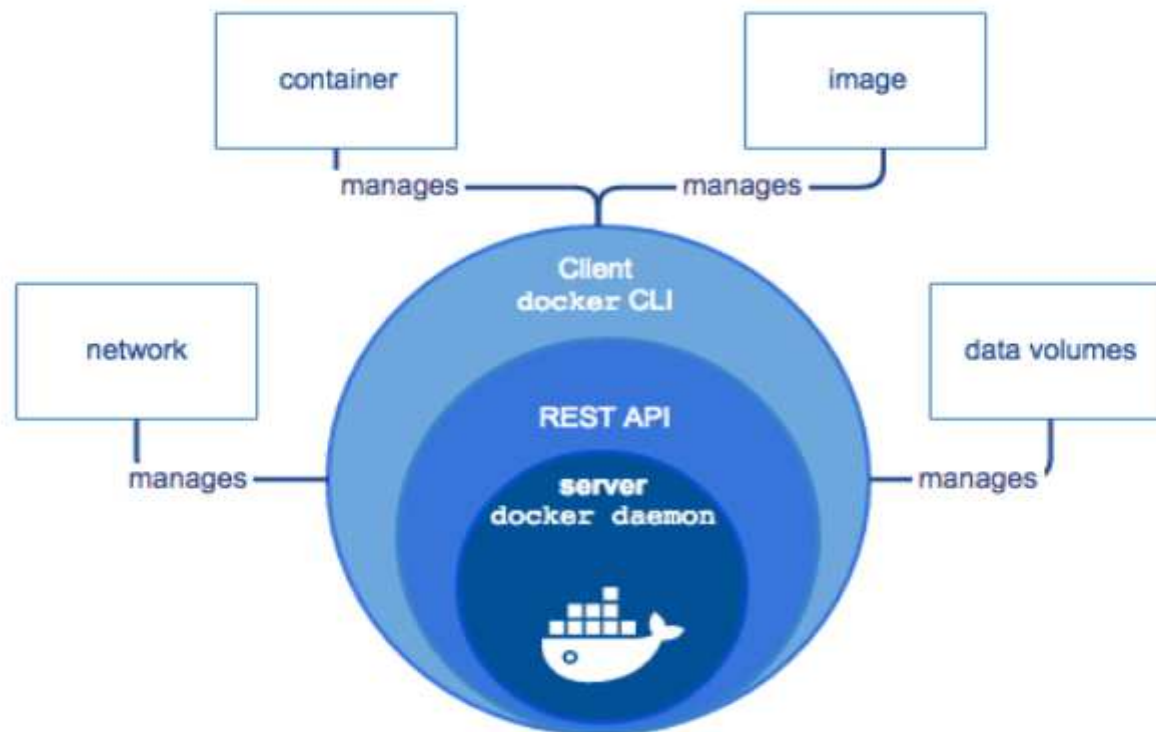


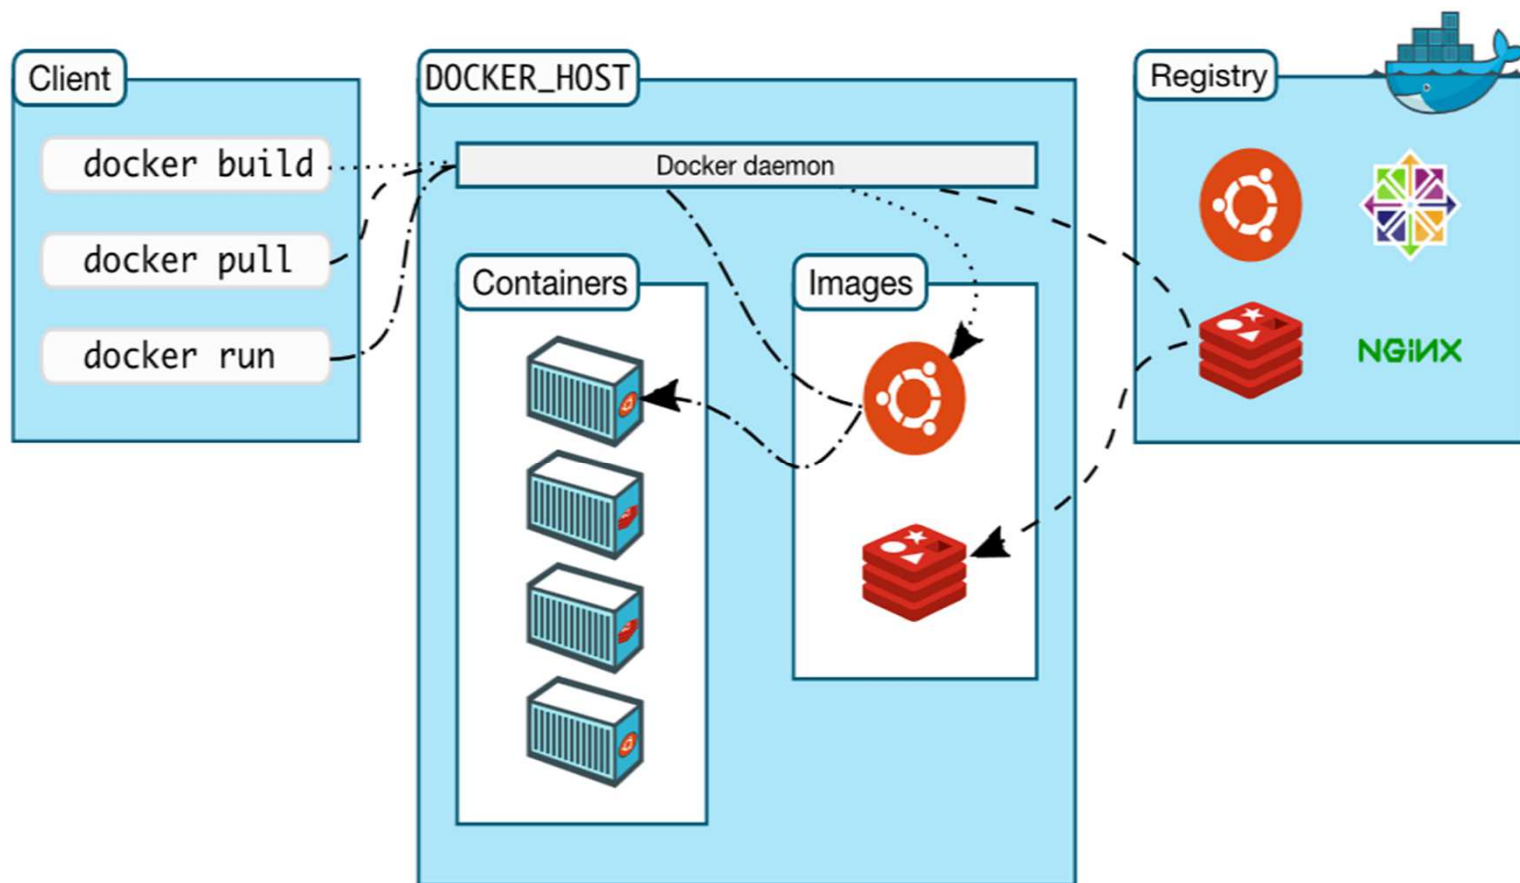
- Kubernetes uses docker as its default container runtime.
- Docker Initial release 13 March 2013
- Operating system Linux, Windows, macOS
- Written in Go
- Original author Solomon Hykes
- Website docker.com

Docker Engine

Docker Engine is a client-server application with these major components:

- A server which is a type of long-running program called a daemon process (the `dockerd` command).
- A REST API which specifies interfaces that programs can use to talk to the daemon and instruct it what to do.
- A command line interface (CLI) client (the `docker` command).





The LifeCycle of a Docker Container I

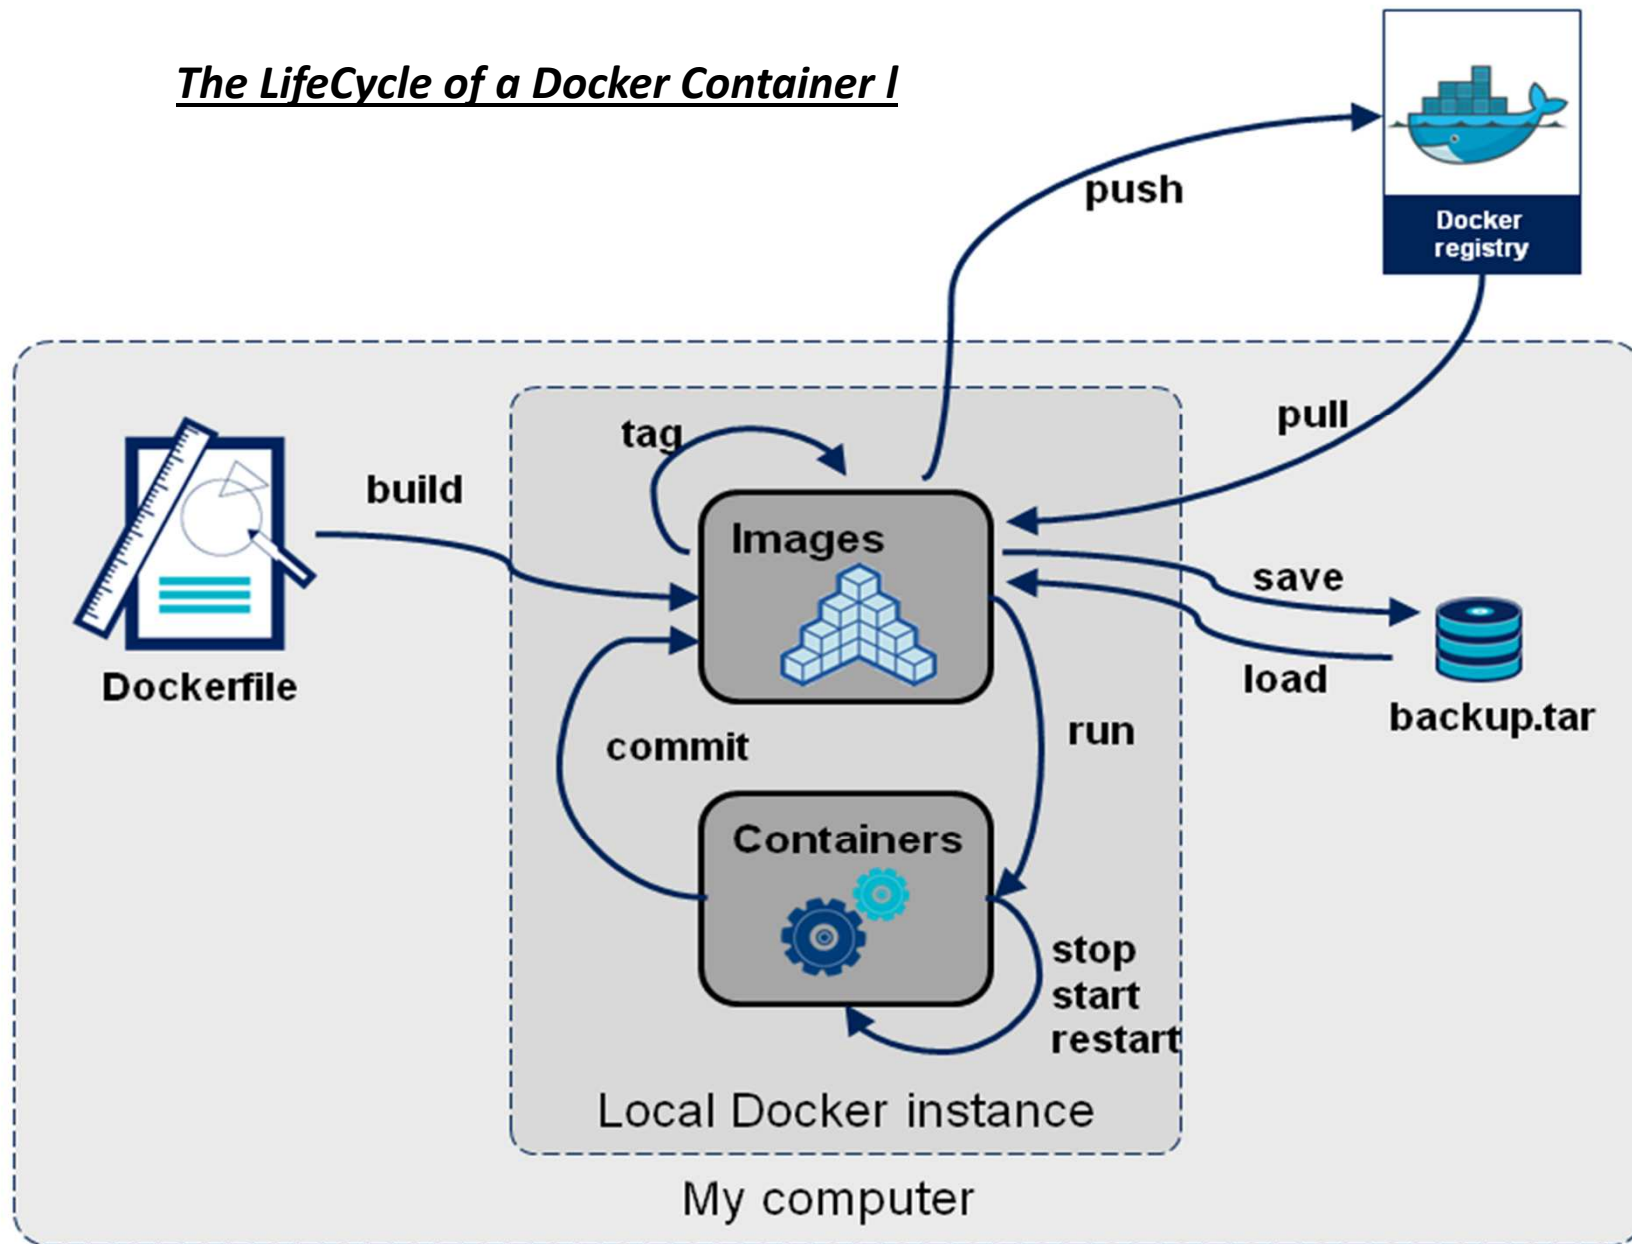


Image Source : <https://stackoverflow.com/questions/23735149/what-is-the-difference-between-a-docker-image-and-a-container>

Docker Editions

Docker Community Edition

Get started with Docker and experimenting with container-based apps. Docker CE is available on many platforms, from desktop to cloud to server. Build and share containers and automate the development pipeline from a single environment. Choose the Edge channel to get access to the latest features, or the Stable channel for more predictability.

[Learn more about Docker CE](#)

Docker Enterprise Edition

Designed for enterprise development and IT teams who build, ship, and run business critical applications in production at scale. Integrated, certified, and supported to provide enterprises with the most secure container platform in the industry to modernize all applications. Docker EE Advanced comes with enterprise [add-ons](#) like UCP and DTR.

[Learn more about Docker EE](#)

Run Docker anywhere



Docker for Mac



Docker for
Windows

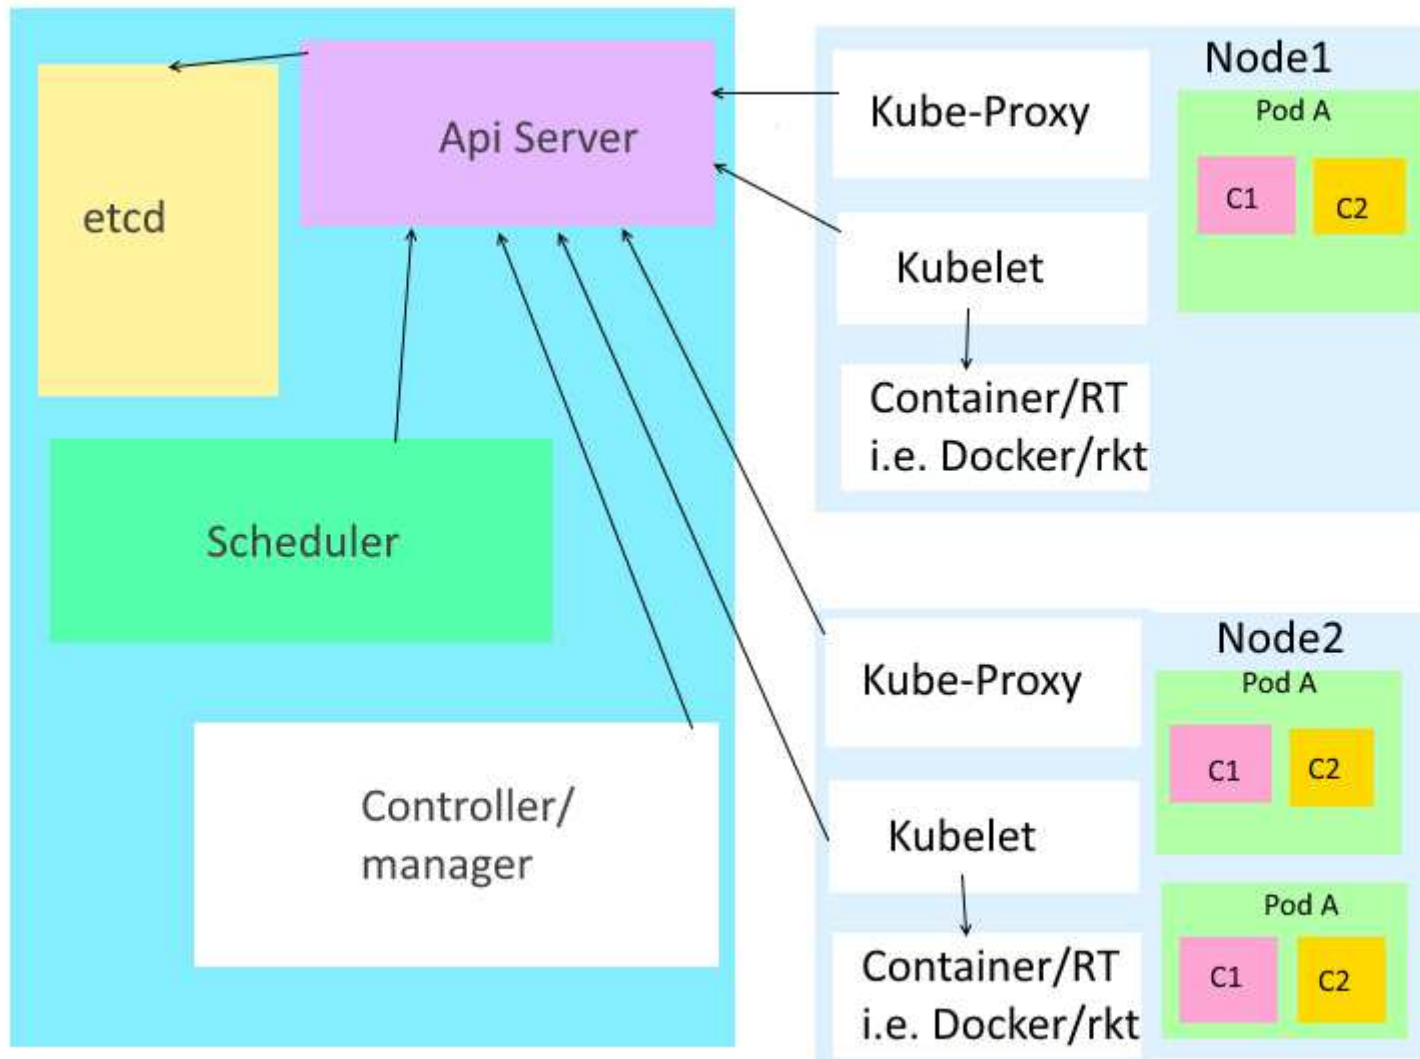


Docker for Linux

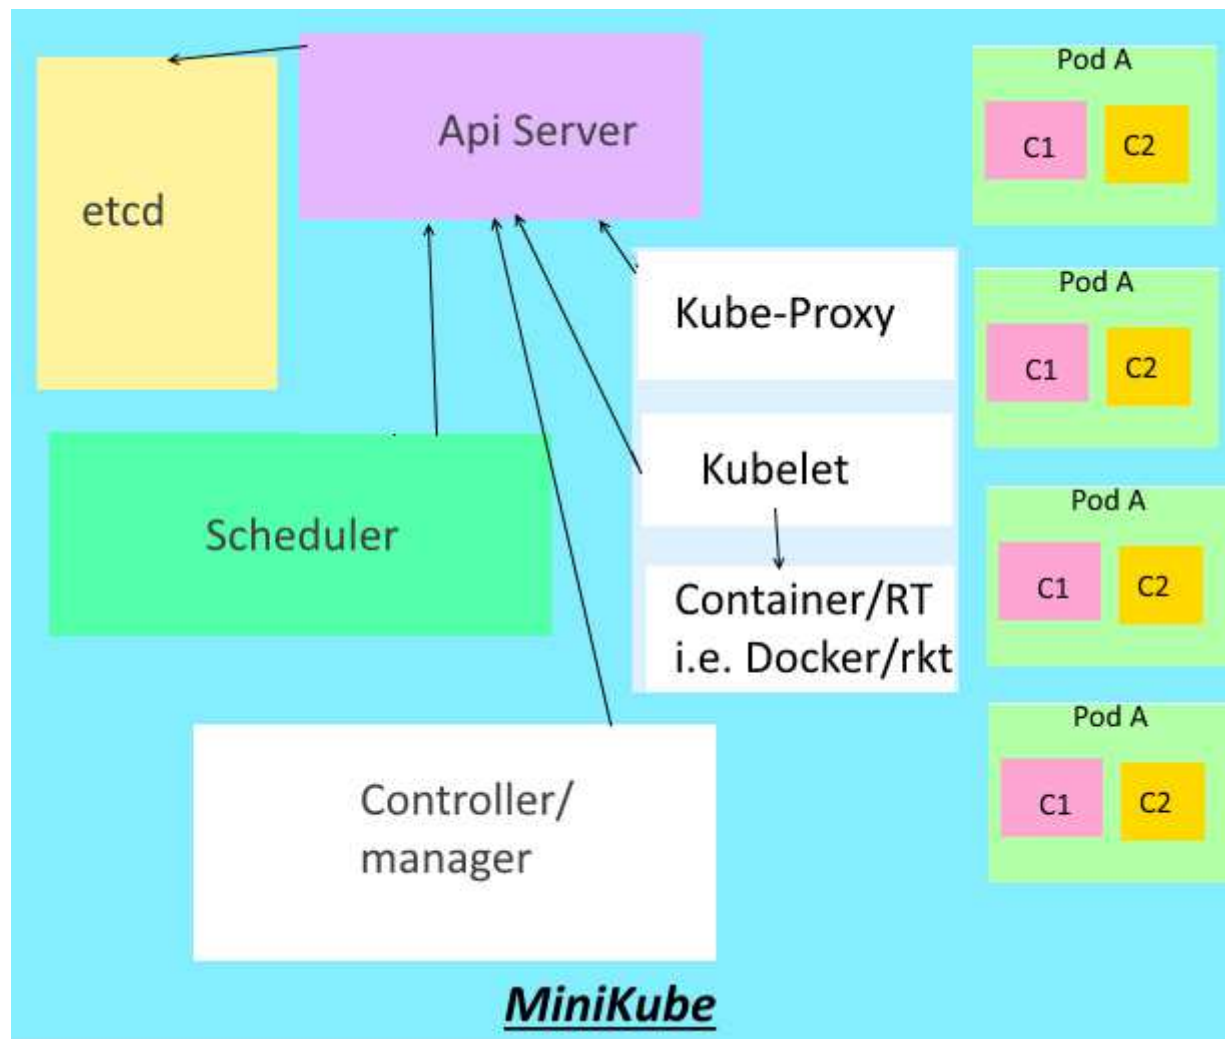
Kubernetes Resources

- NameSpaces
- ReplicationController
- ReplicaSet
- Job
- CronJob
- DaemonSet
- Deployment
- Service, Endpoints , Ingress
- ConfigMap, Secret
- PersistentVolume, PersistentVolumeClaim
- StorageClass

We can use kubectl – the client application to talk to the Api server (an Restful service) or via directly by https. The other way to talk to the Api server is via containers we have deployed in our in Pods. Since the Api server is a restful implementation so can support different versions of a restful Api Service.



Minikube provides us with a Kubernetes Cluster in a Box i.e. a virtual Machine. So we combine the components of a Master and a worker Node. So hence we have no additional nodes. We can use kubectl the client application to create our pods and services the same as with a full cluster(masters & nodes).

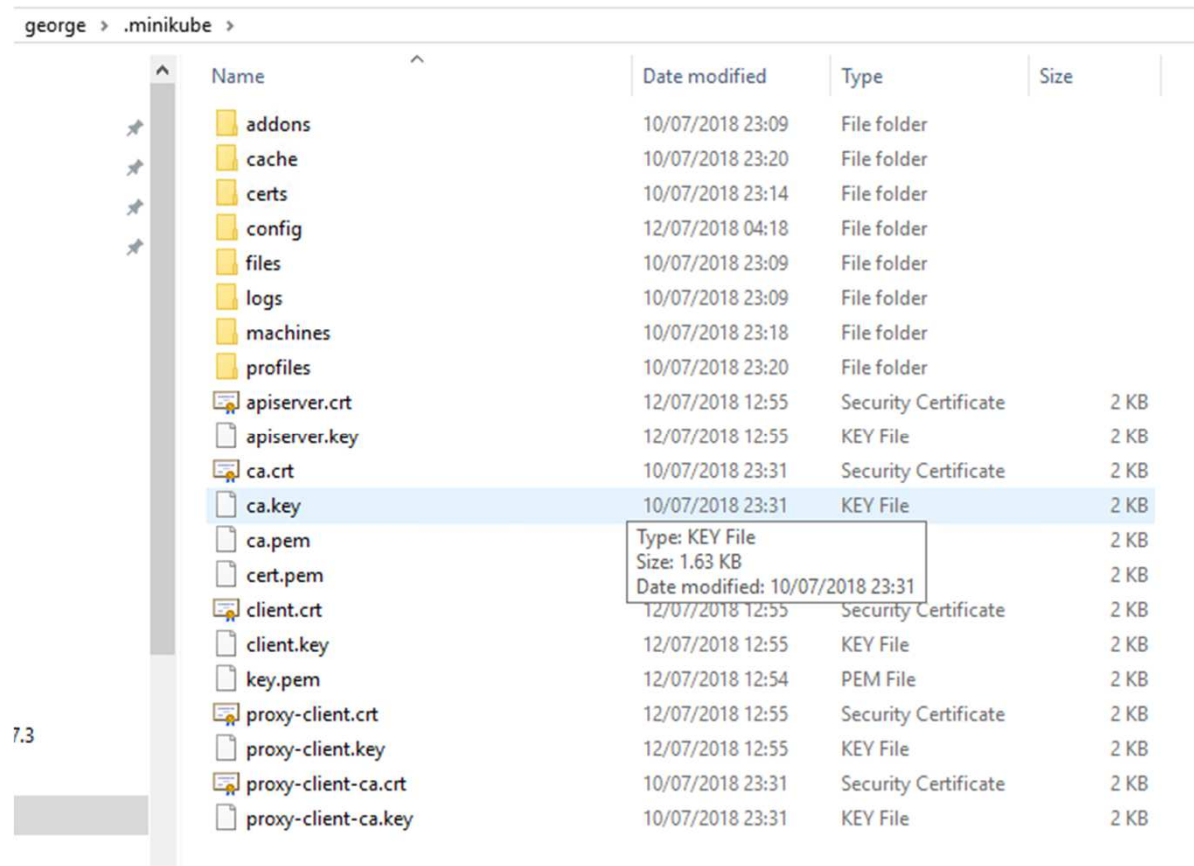


Install/Get Minikube and KubeCtl

- We will install/get minikube and kubectl for windows (hyperv)
- Make a directory i.e. c:\k8\bin
- Next Download minikube-windows-amd64.exe from <https://storage.googleapis.com/minikube/releases/v0.28.0/minikube-windows-amd64.exe>
- Rename minikube-windows-amd64.exe to minikube.exe and move it into c:\k8\bin
- Next download kubectl.exe from <https://storage.googleapis.com/kubernetes-release/release/v1.10.3/bin/windows/amd64/kubectl.exe>
- move minikube.exe into c:\k8\bin
- Next add c:\k8\bin to your path (please note docker also now bundles kubectl.exe in his bin folder so decide if you want to use that or use you own. Put your bin folder first in the path if you want to make sure of using your own.
- by default MiniKube install into your home directory

MiniKube default home directory

- by default MiniKube installs into your home directory under directory .minikube

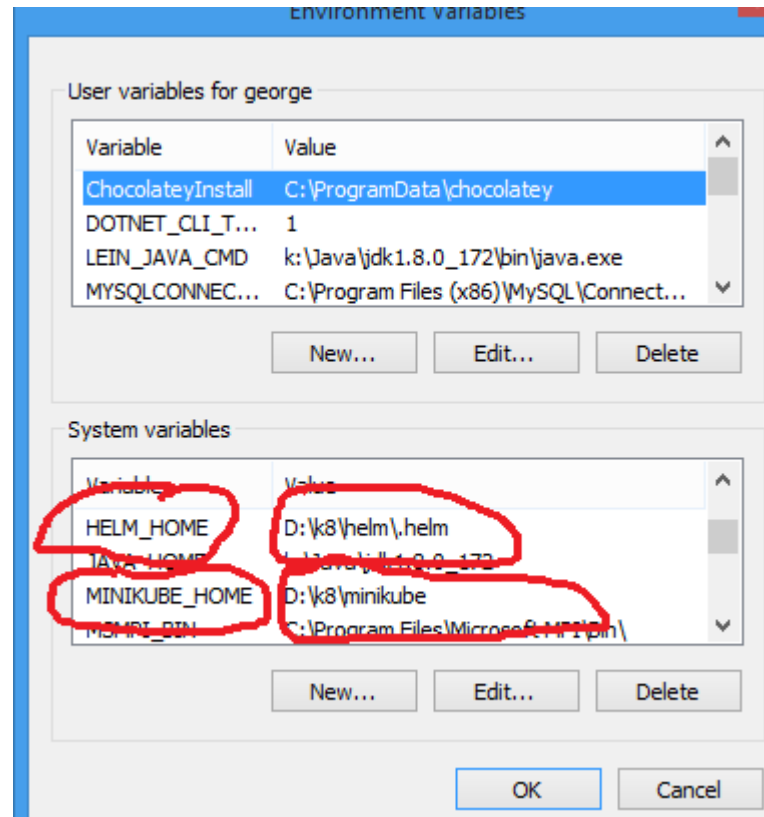


Name	Date modified	Type	Size
addons	10/07/2018 23:09	File folder	
cache	10/07/2018 23:20	File folder	
certs	10/07/2018 23:14	File folder	
config	12/07/2018 04:18	File folder	
files	10/07/2018 23:09	File folder	
logs	10/07/2018 23:09	File folder	
machines	10/07/2018 23:18	File folder	
profiles	10/07/2018 23:20	File folder	
apiserver.crt	12/07/2018 12:55	Security Certificate	2 KB
apiserver.key	12/07/2018 12:55	KEY File	2 KB
ca.crt	10/07/2018 23:31	Security Certificate	2 KB
ca.key	10/07/2018 23:31	KEY File	2 KB
ca.pem			2 KB
cert.pem			2 KB
client.crt	12/07/2018 12:55	Security Certificate	2 KB
client.key	12/07/2018 12:55	KEY File	2 KB
key.pem	12/07/2018 12:54	PEM File	2 KB
proxy-client.crt	12/07/2018 12:55	Security Certificate	2 KB
proxy-client.key	12/07/2018 12:55	KEY File	2 KB
proxy-client-ca.crt	10/07/2018 23:31	Security Certificate	2 KB
proxy-client-ca.key	10/07/2018 23:31	KEY File	2 KB

But you can set a windows environment variable to point to another location i.e. if you have limited space on your c disk or want the vm that MiniKube creates to run another disk for performance.

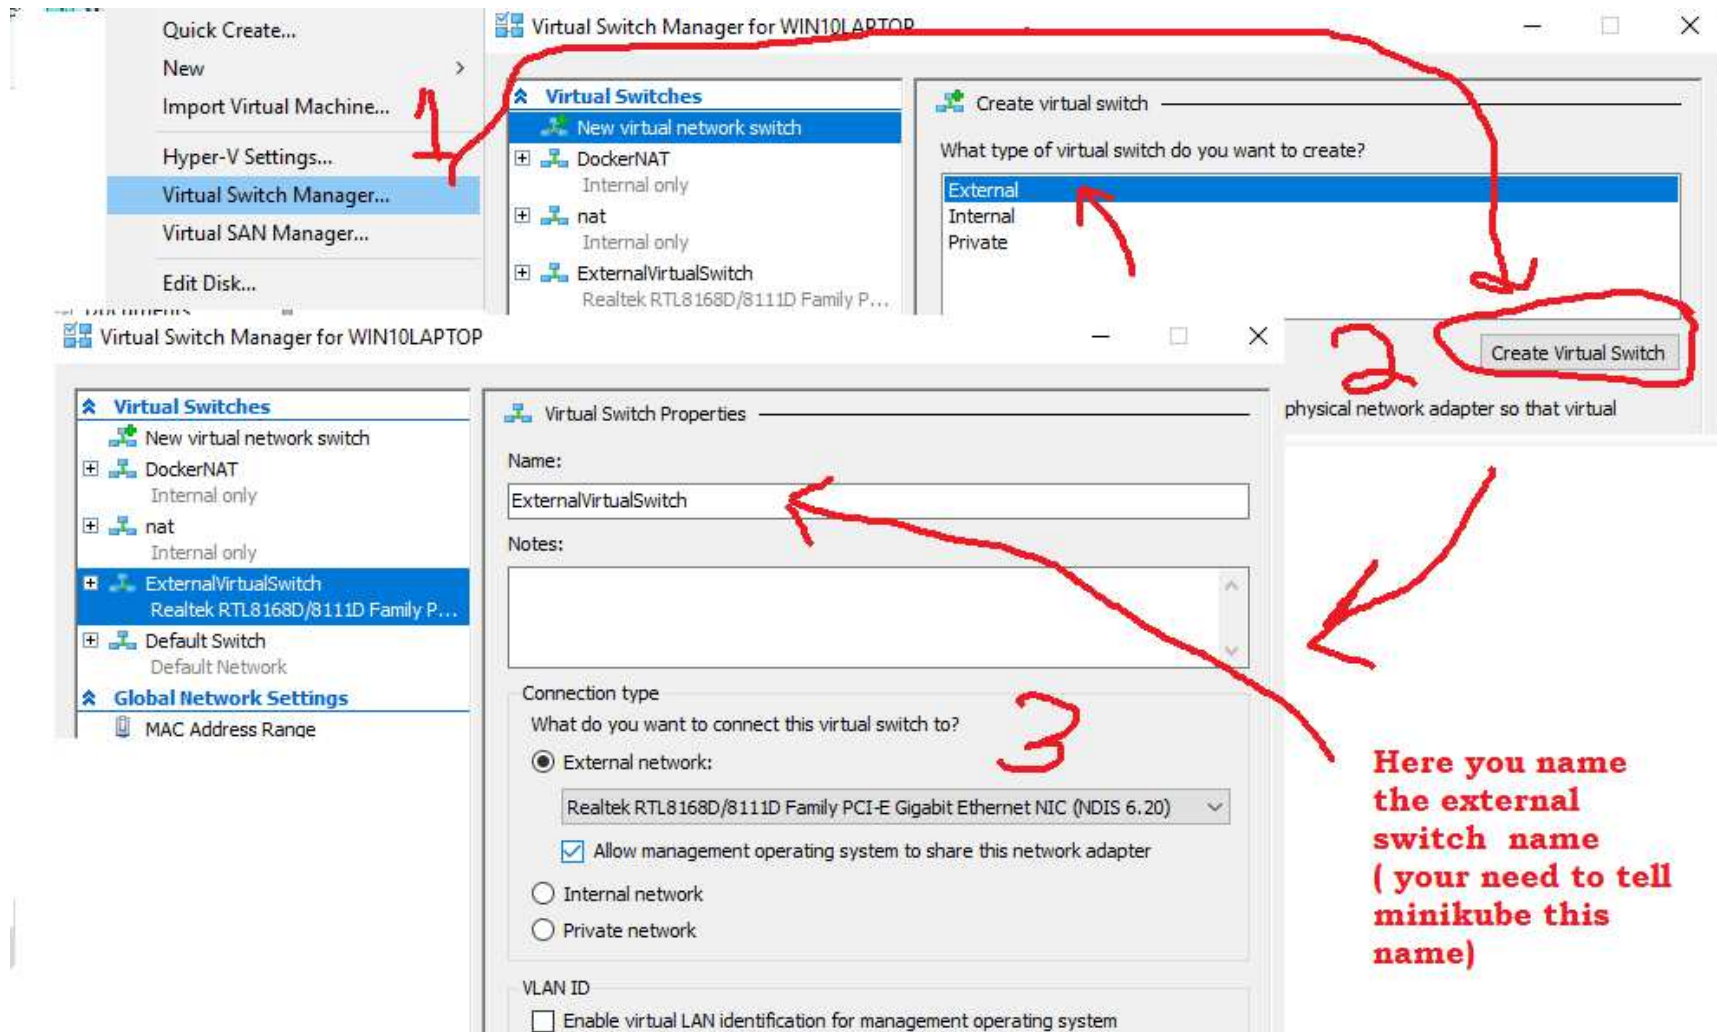
Changing the MiniKube home directory

- Set MINIKUBE_HOME to your new location – the install process will place it there instead



You will also notice you can set the Helm home directory from the default of your user home directory, as in this case we moved it to D:\k8\helm\.helm. Helm is a kubernetes packages manager, we haven't discussed that yet. But just for reference encase you need to.

How to install MiniKube on HyperV (First create a Hyperv Virtual External Switch)



By default Minikube default to creating a Virtualbox VM as it default. If we want to use HyperV we must first create a external virtual switch in Hyperv (see graphics) you have have spaces or _ but best to keep it simple unlike my example !!!

How to install MiniKube on Hyperv

```
C:\ Windows Command Processor - minikube start --vm-driver hyperv --hyperv-virtual-switch "External_Virtual_Switch"
Microsoft Windows [Version 10.0.17134.112]
(c) 2018 Microsoft Corporation. All rights reserved.

Convert Virtual Disk
 40%
 [ooooooooooooooooooooooooooooooooooooooooooooooooooooooooooooo]

C:\Windows\System32>d:

D:\>cd d:\k8\bin

d:\k8\bin>minikube start --vm-driver hyperv --hyperv-virtual-switch "External_Virtual_Switch"
Starting local Kubernetes v1.10.0 cluster...
Starting VM...
Downloading Minikube ISO
 153.08 MB / 153.08 MB [=====] 100.00% 0s
```

We tell Minikube to use hyperv to create our vm machine and tell it the external virtual switch to use. (note my virtual switch has underscores in it, that's because I did several installations on different machines and named them different names that's all).

How Stop and Start MiniKube on Hyperv

Shutting Down

0%

[

```
D:\k8\bin>minikube stop
Stopping local Kubernetes cluster...
```

Administrator: Windows PowerShell

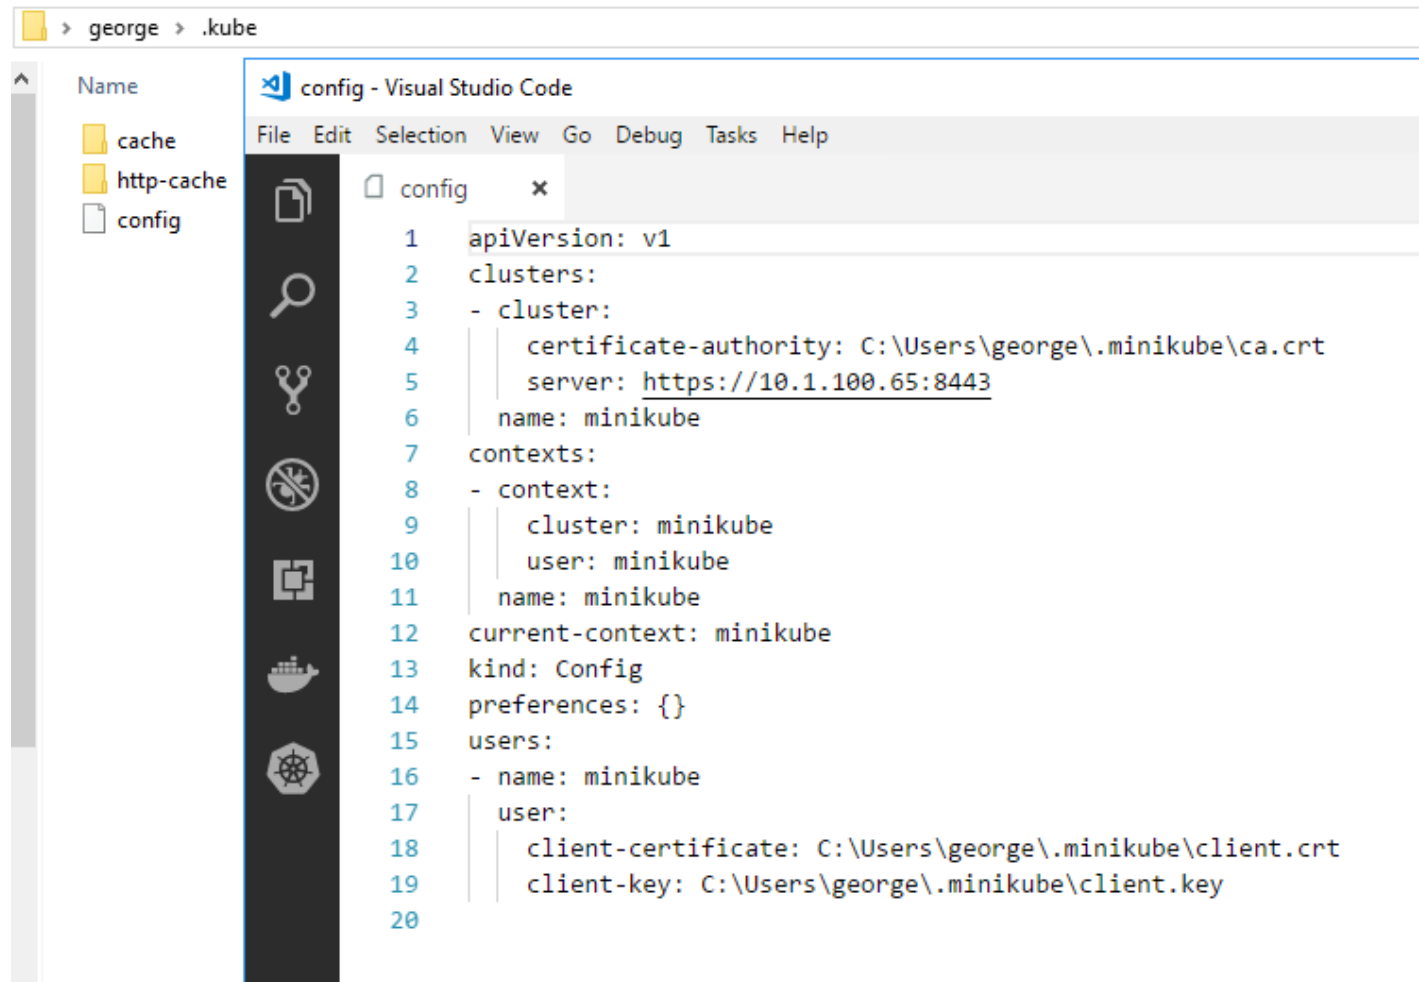
```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

PS C:\WINDOWS\system32> minikube start
Starting local Kubernetes v1.10.0 cluster...
Starting VM...
Getting VM IP address...
Moving files into cluster...
Setting up certs...
Connecting to cluster...
Setting up kubeconfig...
Starting cluster components...
Kubectl is now configured to use the cluster.
Loading cached images from config file.
```

Be warned this never seems to work without errors in the current version of 0.28 I was using. After we have installed Minikube and it is up and running we can just use `minikube start`, as it creates a config file in your home directory under `.kube`

Minikube Creates a Cluster Config file

After we have installed Minikube , Minikube creates a config file in your home directory under .kube called config. Lets have a look at it.

A screenshot of the Visual Studio Code editor. The left sidebar shows a file explorer with the path 'george > .kube' and three files: 'cache', 'http-cache', and 'config'. The 'config' file is selected. The main editor window shows the content of the 'config' file, which is a YAML configuration for the Minikube cluster. The file is named 'config' and has a tab icon with a close button. The content is as follows:

```
1  apiVersion: v1
2  clusters:
3  - cluster:
4      certificate-authority: C:\Users\george\.minikube\ca.crt
5      server: https://10.1.100.65:8443
6      name: minikube
7  contexts:
8  - context:
9      cluster: minikube
10     user: minikube
11     name: minikube
12  current-context: minikube
13  kind: Config
14  preferences: {}
15  users:
16  - name: minikube
17     user:
18         client-certificate: C:\Users\george\.minikube\client.crt
19         client-key: C:\Users\george\.minikube\client.key
20
```

The Purpose of this is so that client tools that need to work with the cluster have a way to find out how to connect to the cluster and hence the kubectl client tool can send your command to the Kubernetes clusters Api server. (we can cluster context with `kubectl config set-context`)

MiniKube addons

Minikube has the option to enable/disable various addons – there are issues around addons and you need to make sure the various configuration files are updated. i.e. perhaps needing to restart Minikube. But this is just a passing note to make you aware of the addon feature.

```
C:\Windows\System32>minikube addons list
- addon-manager: enabled
- coredns: disabled
- dashboard: enabled
- default-storageclass: enabled
- efk: disabled
- freshpod: disabled
- heapster: disabled
- ingress: disabled
- kube-dns: enabled
- metrics-server: disabled
- registry: disabled
- registry-creds: disabled
- storage-provisioner: enabled
```


Kubernetes namespaces

A Kubernetes cluster has namespaces lets have a look and use the kubectl to talk the api server and send some commands.

```
D:\k8\bin> kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	etcd-minikube	1/1	Running	0	12m
kube-system	kube-addon-manager-minikube	1/1	Running	1	1h
kube-system	kube-apiserver-minikube	1/1	Running	0	12m
kube-system	kube-controller-manager-minikube	1/1	Running	0	12m
kube-system	kube-dns-86f4d74b45-hq45c	3/3	Running	1	1h
kube-system	kube-proxy-vflz5	1/1	Running	0	12m
kube-system	kube-scheduler-minikube	1/1	Running	0	1h
kube-system	kubernetes-dashboard-5498ccf677-pcth4	1/1	Running	4	1h
kube-system	registry-h6xcw	1/1	Running	0	28m
kube-system	storage-provisioner	1/1	Running	2	1h

```
D:\k8\bin> kubectl get services --all-namespaces
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1h
kube-system	kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP	1h
kube-system	kubernetes-dashboard	NodePort	10.110.92.97	<none>	80:30000/TCP	1h
kube-system	registry	ClusterIP	10.98.77.99	<none>	80/TCP	28m

```
D:\k8\bin> kubectl get rc --all-namespaces
```

NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE
kube-system	registry	1	1	1	29m

```
D:\k8\bin> kubectl get pods
```

No resources found.

Lets ssh into the MiniKube vm

Use minikube ssh

To ssh into your minikube – be warned if you have the windows 10 ssh comand on the path that won't work. I use git's ssh and put that on the path .

```
PS C:\WINDOWS\system32> minikube ssh
```



```
/bin /boot /dev /etc /home /lib /lib64 /media /mnt /opt /sbin /usr /var
```

```
$ ls /etc/docker/  
ca.pem certs.d key.json server-key.pem server.pem  
$ -
```

I created this directory for my self - signed certs. But you see, we have access to the Minikube VM and change stuff via su or sudo. But NOTE: only some directories are persisted.

Lets run Statefulset of poker game tables

I will talk through the various stages and will demo this if there is time.

```
PS C:\k8\bin> .\kubect1.exe delete sts poker
statefulset.apps "poker" deleted
PS C:\k8\bin> .\kubect1.exe get pods
No resources found.
PS C:\k8\bin> .\kubect1.exe create -f files/simplepodexamples/pokerstatefulset.yaml
statefulset.apps "poker" created
PS C:\k8\bin> .\kubect1.exe get pods
NAME      READY   STATUS    RESTARTS   AGE
poker-0   1/1     Running   0           4s
poker-1   1/1     Running   0           2s
poker-2   1/1     Running   0           0s
PS C:\k8\bin> .\kubect1.exe get sts
NAME      DESIRED   CURRENT   AGE
poker     3         3         2h
PS C:\k8\bin> .\kubect1.exe get sc
NAME      PROVISIONER      AGE
standard (default)  k8s.io/minikube-hostpath  1d
PS C:\k8\bin> .\kubect1.exe get rs
No resources found.
PS C:\k8\bin> .\kubect1.exe get rc
No resources found.
PS C:\k8\bin> .\kubect1.exe get jobs
No resources found.
PS C:\k8\bin> .\kubect1.exe get cronjobs
No resources found.
```

Diagram illustrating the sequence of commands and their outputs, with red arrows and numbers indicating the flow:

1. Arrow pointing to the command: `.\kubect1.exe delete sts poker`
2. Arrow pointing to the command: `.\kubect1.exe get pods`
3. Arrow pointing to the command: `.\kubect1.exe create -f files/simplepodexamples/pokerstatefulset.yaml`
4. Arrow pointing to the command: `.\kubect1.exe get pods`
5. Arrow pointing to the command: `.\kubect1.exe get sts`
6. Arrow pointing to the command: `.\kubect1.exe get sc`
7. Arrow pointing to the command: `.\kubect1.exe get rs`
8. Arrow pointing to the command: `.\kubect1.exe get rc`
9. Arrow pointing to the command: `.\kubect1.exe get jobs`
10. Arrow pointing to the command: `.\kubect1.exe get cronjobs`

Lets run Statefulset of poker game tables

Lets talk Yaml – The yaml layout is very important, with regards to space and indentation. Here we run a Statefulset , which controls a service with 3 replicas. This allows our game tables to have a consistent identify, so if an pod gets restarted then the pod gets the same identify. Also I am using my own local registry with a self signed cert – for the image cpucode.local/pokerv1.

</> pokerstatefulset.yaml

```
1  apiVersion: apps/v1beta1
2  kind: StatefulSet
3  metadata:
4    name: poker
5  spec:
6    serviceName: poker
7    replicas: 3
8    template:
9      metadata:
10       labels:
11         app: poker
12       spec:
13         hostAliases:
14           - ip: "10.1.100.133"
15             hostnames:
16               - "cpucode.local"
17           - ip: "10.1.100.150"
18             hostnames:
19               - "anotherhost.local"
20               - "andanotherhost.local"
21         containers:
22           - name: pokerv1
23             image: cpucode.local/pokerv1
24
25   # ports:
26   #   - name: http
27   #     containerPort: 8080
28   # volumeMounts:
29   #   - name: data
30   #     mountPath: /var/data
31   # volumeClaimTemplates:
32   #   - metadata:
33   #     name: data
34   #   spec:
35   #     resources:
36   #       requests:
37   #         storage: 1Mi
38   #     accessModes:
39   #       - ReadWriteOnce
```

Lets run a single poker game table pod

When we run a single pod that is not controlled by a controller, then we don't get any automatic behaviour to restart that pod if i.e. the node that the pod was running on in the cluster fails.

Basically we want to control our pods via services and those services by controllers. So that we can scale up and down the resources as need. In terms of microservices we may want to do blue/green deployment. We have a Deployment controller that does actually that for us.

The image shows a code editor on the left with a YAML file named `pokerstatefulset.yaml` and a terminal on the right showing the execution of `kubectrl` commands and the resulting pod status.

YAML File Content:

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: poker-pod
5  spec:
6    hostAliases:
7      - ip: "10.1.100.133"
8        hostnames:
9          - "cpucode.local"
10     - ip: "10.1.100.150"
11       hostnames:
12         - "anotherhost.local"
13         - "andanotherhost.local"
14
15   containers:
16     - image: cpucode.local/pokerv1
17       name: pokerv1
```

Terminal Output:

```
PS C:\k8\bin> .\kubectrl.exe run singlepokertable --image=cpucode.local/pokerv1
deployment.apps "singlepokertable" created
PS C:\k8\bin> .\kubectrl.exe get pods
```

NAME	READY	STATUS	RESTARTS	AGE
poker-0	1/1	Running	0	3h
poker-1	1/1	Running	0	3h
poker-2	1/1	Running	0	3h
singlepokertable-5dc677fbcd-29vvg	1/1	Running	0	21m

PS C:\k8\bin> █

Annotations:

- A black arrow points from the `cpucode.local/pokerv1` image name in the terminal output to the `image: cpucode.local/pokerv1` line in the YAML file.
- A black arrow points from the `singlepokertable-5dc677fbcd-29vvg` pod name in the terminal output to the `name: poker-pod` line in the YAML file.
- Text below the terminal: "Here we use the " kubectrl run " to run the pokertable pod. We could have used "kubectrl create -f filename" to run it but just demo the run option."
- Text below the terminal: "Here we run a single pod see (Kind: pod)"

Installing Helm into our Cluster



A terminal window showing the installation of Helm. The commands and output are as follows:

```
D:\k8\bin>echo %HELM_HOME%
D:\k8\helm\.helm

D:\k8\bin>helm init
Creating D:\k8\helm\.helm\repository
Creating D:\k8\helm\.helm\repository\cache
Creating D:\k8\helm\.helm\repository\local
Creating D:\k8\helm\.helm\plugins
Creating D:\k8\helm\.helm\starters
Creating D:\k8\helm\.helm\cache\archive
Creating D:\k8\helm\.helm\repository\repositories.yaml
Adding stable repo with URL: https://kubernetes-charts.storage.googleapis.com
Adding local repo with URL: http://127.0.0.1:8879/charts
$HELM_HOME has been configured at D:\k8\helm\.helm.
Warning: Tiller is already installed in the cluster.
(Use --client-only to suppress this message, or --upgrade to upgrade Tiller to the current version.)
Happy Helming!

D:\k8\bin>
```

Yellow annotations are present in the image:

- A yellow arrow points from the text "1" to the command `helm init`.
- A yellow arrow points from the text "2" to the output line `Creating D:\k8\helm\.helm\cache\archive`.
- A yellow arrow points from the text "3" to the output line `$HELM_HOME has been configured at D:\k8\helm\.helm.`

Helm has the client side which and a server side called Tiller that get installed into your cluster. As noted above you use (`helm init`). You can now install Helm charts (i.e. names to do with ships) charts are packages of Yaml files so you can deploy large or complex applications i.e. MongoDB when you want to install a MongoDB replicaset etc.

The End for Now

- This was just a 15 minute lightning talk with Demo of a scaling a poker game. The talk was more focused on installing and setting up Minikube. I would have loved to discussed deploying microservices via the different types of controllers in Kuterbetes. But that is a different talk for another time perhaps.
- Thanks – See you next time at the
Limerick Ai Software Development meetup